

QECTOR Decoder Workbench

USER MANUAL

Linux Edition

Application version 0.5.2

Decoder backend qector-decoder-v3 0.7.0 (bundled wheel, activated offline on first launch; minimum supported 0.7.0)

56-tool MCP server, 16 decoders, 10 code families

Guillaume Lessard / iD01t Productions. <https://www.qector.store>

Generated 2026-08-02

Contents

1. Introduction	3
2. System requirements	3
3. Installation	3
Ubuntu, Debian, Mint	3
antiX, MX and other Debian-family distributions	3
4. First launch	3
5. The graphical interface	4
6. Working with codes	5
7. Decoders	5
8. Decoding, benchmarking and streaming	6
Decoder Lab	6
Benchmark	6
Batch and Streaming	6
9. Hardware, diagnostics and export	6
Hardware	6
Diagnostics	7
Documentation export	7
10. The MCP server	8
11. Environment variables	8
12. Troubleshooting	8
13. Licensing and support	9
14. Glossary	9
Appendix A. Full MCP tool list	10

1. Introduction

QECTOR Decoder Workbench is a professional desktop suite for building quantum error-correcting codes and analysing their decoders. It combines a graphical workbench of seven feature tabs and a live console, a resilient decoding backend, a multi-format documentation generator, and a headless server that exposes 56 tools over the Model Context Protocol (MCP) for use by automation and language-model agents.

The suite is powered by `qector-decoder-v3`, a compiled Rust decoding engine with a Python API. It offers 16 single-shot decoders across 10 code families, batch and streaming decode, hardware detection, and fully reproducible seeded runs.

Note. This edition is fully self-contained: the decoder backend wheel is bundled and activates offline on first launch. Reported logical error rate, throughput and latency are simulation outputs that depend on the seed, the hardware and the workload; regenerate them for your own setting before quoting.

2. System requirements

- A 64-bit x86_64 Linux desktop built on glibc 2.31 or newer: Ubuntu 20.04+, Debian 11+, Linux Mint 20+, antiX 21+, MX 21+, Fedora 33+, openSUSE Leap 15.3+, recent Arch, and anything newer.
- About 150 MB of free disk space.
- The decoder wheel is bundled; it activates automatically on first launch, entirely offline (no internet, pip, or extra Python needed).

3. Installation

Ubuntu, Debian, Mint

- `sudo dpkg -i ./qector-workbench_0.5.2_amd64.deb`
- `sudo apt-get -f install` (only if dpkg reports missing dependencies)
- A QECTOR icon appears in the Science menu.
- Remove it later with: `sudo apt remove qector-workbench`

antiX, MX and other Debian-family distributions

- `sudo dpkg -i ./qector-workbench_0.5.2_amd64.deb`
- `sudo apt-get -f install` (only if dpkg reports missing dependencies)

4. First launch

The application starts immediately. The bundled decoder backend wheel is extracted and activated on first launch, no internet access is needed at any point. Any outdated managed decoder left by an older release is purged automatically first. The window title and the status bar show the workbench and decoder versions.

Start the graphical application with qector-workbench. To run the headless MCP server, which needs no display, use:

```
qector-workbench --mcp
```

Note. To move to a newer decoder, install a newer release of the application; each release bundles its matching decoder wheel.

5. The graphical interface

The workbench presents the following tabs. A typical session builds a code in the Code Explorer, then decodes, benchmarks or streams it in the other tabs.

Code Explorer	Pick one of the nine code families, set the distance, and inspect the code: qubit and check counts, code distance, and a readable Tanner graph. This is where you build the code that the other tabs operate on.
Decoder Lab	Run a single seeded decode with any of the ten decoders. The panel reports the correction weight, whether the correction reproduces the observed syndrome (syndrome_valid), and whether a logical operator was flipped. If a chosen decoder cannot handle the selected code, a resilient fallback recovers automatically and records the full attempt trace.
Benchmark	Sweep a decoder over many seeded shots and read throughput (decodes per second), latency percentiles, and the logical error rate for that workload.
Batch and Streaming	Decode a batch of syndromes on the CPU, CUDA or OpenCL backend, and run a sliding-window streaming session with per-round commit telemetry.
Hardware	Detect CUDA and OpenCL availability, show system information, and read a decoder recommendation that never suggests a decoder the selected code cannot use.
Diagnostics	Run a full self-test of the environment, probe which decoders work for a given code, and exercise the resilient decode path with a complete trace.
Documentation	Export a provenance-stamped record of the current code in six formats: Markdown, JSON, HTML, LaTeX, PDF and SVG.
Console	A live, severity-coloured log of every operation, kept in one place for review and troubleshooting.

6. Working with codes

Open the Code Explorer, choose a family, set the distance (minimum 3), and build. The 10 families are:

repetition	Repetition Code	graphlike	1D chain that protects against bit-flip errors. The simplest teaching code.
ring	Ring Code	graphlike	Periodic 1D code (repetition on a ring). Good for streaming and cyclic examples.
rotated_surface	Rotated Surface Code	graphlike	Compact rotated surface code; the standard workhorse for 2D QEC studies.
unrotated_surface	Unrotated Surface Code	graphlike	Classic surface code layout with boundary stabilisers.
toric	Toric Code	graphlike	Surface code on a torus (periodic boundaries); two logical qubits.
heavy_hex	Heavy Hex Code	graphlike	IBM heavy-hexagon layout used by superconducting hardware.
bicycle	Bicycle Code (qLDPC)	qLDPC (non graphlike)	A qLDPC code that every applicable decoder can handle.
bivariate_bicycle	Bivariate Bicycle Code (qLDPC)	qLDPC (non graphlike)	The IBM bivariate-bicycle family, e.g. the $[[72,12,6]]$ gross code.
hypergraph_product	Hypergraph-Product Code (CSS)	graphlike	Constructs a CSS code from two classical codes; graphlike here.
color_code	Color Code (Triangular 4.8.8)	graphlike	

Note. The two qLDPC families (bicycle, bivariate_bicycle) are not graphlike. Decode them with bp_osd, sparse_blossom, blossom, hybrid, predecoded, auto or auto_router. The resilient decode path selects a working decoder automatically if you pick one that does not apply.

7. Decoders

These single-shot decoders are available by exact name in the Decoder Lab, Benchmark and Diagnostics tabs and in the MCP server.

union_find	Fast graphlike matching for surface, repetition and ring codes. Not for qLDPC.
fast_union_find	SIMD-accelerated union-find; highest throughput on graphlike codes. Not for qLDPC.
blossom	Exact minimum-weight perfect matching. Best accuracy on graphlike codes; slower.
sparse_blossom	Region-growing matcher; strong accuracy with better scaling than exact blossom.
bp_osd	Belief propagation with ordered-statistics decoding. The go-to decoder for qLDPC codes.
auto	Policy decoder: inspects the code and dispatches a suitable concrete decoder automatically.
hybrid	Neural pre-decoder feeding sparse blossom; adaptive edge weights for enriched decoding.

lookup_table	Exact table for small codes only (refused above 20 checks). Constant-time lookups.
predecoded	A pre-decoding stage in front of a matcher to reduce work on likely error patterns.
auto_router	Routes graphlike codes to matching and qLDPC codes to BP-OSD; safe default everywhere.
hybrid_cascade	Hybrid Cascade , Union-Find pre-filter + Blossom/BP-OSD escalation : trivial syndromes resolve at UF speed, hard ones escalate to the accurate decoder. Exposes prefilter_hits / escalations / prefilter_hit_rate stats. Graphlike codes only (UF pre-filter).
gnn_belief_matching	GNN Belief Matching , GNN-predicted per-qubit weights guide a weighted matching decode (GNNBeliefMatcher); a built-in faithfulness check falls back to plain MWPM so corrections stay syndrome-valid. Graphlike codes.
belief_matching	Belief Matching , sum-product BP posteriors reweight an exact Blossom matching step (Higgott et al. 2023), recovering error-correlation information plain MWPM discards. Faithfulness-checked; falls back to plain MWPM so corrections stay syndrome-valid.
two_stage	Two-Stage , decoupled X and Z sector decoders for CSS / color codes; decodes X and Z check sub-graphs separately with configurable base decoders.
ambiguity_cluster	Ambiguity Cluster , cluster-growth decoder for high noise or non-graphlike syndromes; partitions ambiguous checks into local clusters before solving.
colour_code	Colour Code , BP-OSD hypergraph decoder over undecomposed detector error models (v0.7.0); preserves multi-detector 3-body error mechanisms that standard MWPM discards.

8. Decoding, benchmarking and streaming

Decoder Lab

Select a decoder, set a physical error rate and a seed, and run a single decode. The result reports the correction weight, syndrome_valid (whether the correction reproduces the syndrome), and logical_failure (whether a logical operator flipped). A seed makes the run exactly reproducible.

Benchmark

Sweep the chosen decoder over many seeded shots to read throughput (decodes per second), latency percentiles, and the logical error rate for that workload.

Batch and Streaming

Decode a batch of syndromes on the CPU, CUDA or OpenCL backend, or run a sliding-window streaming session with per-round commit telemetry. The CPU backend is the default and works everywhere; GPU backends are used only when present and healthy.

9. Hardware, diagnostics and export

Hardware

Detects CUDA and OpenCL availability and system information, and gives a decoder recommendation that never proposes a decoder the selected code cannot use.

Diagnostics

Runs a full environment self-test, probes which decoders work for a given code, and exercises the resilient decode path with a complete attempt trace. Use it first when anything looks wrong.

Documentation export

The Documentation tab writes a provenance-stamped record of the current code in six formats (Markdown, JSON, HTML, LaTeX, PDF, SVG) into your data directory. For a very large code the SVG export falls back to a minimal valid file rather than failing.

10. The MCP server

The application ships a server that speaks newline-delimited JSON-RPC 2.0 over standard input and output, protocol version 2024-11-05. Start it with the `--mcp` flag. A client sends an initialize request, then an initialized notification, then calls `tools/list` to enumerate the 56 tools or `tools/call` to invoke one.

```
qector-workbench --mcp
```

A complete developer walkthrough is in the QECTOR MCP Integration Guide, and a machine-readable description of every tool is in `QECTOR_LLM_Manual.json`.

11. Environment variables

QECTOR_DATA_DIR	Relocate all per-user data (logs, exports, settings) to a chosen directory.
QECTOR_PYTHON	Legacy override: path to a compatible CPython used only for source-run installs. The shipped app never needs it, the decoder comes from the bundled wheel, offline.
QECTOR_DISABLE_OPENCL	Set to 1 to skip OpenCL probing on systems with unstable GPU drivers. It cannot enable OpenCL; the published wheel is built without that backend.
QECTOR_SILENT	Set to 1 to suppress the backend startup notice.

The per-user data directory for this edition is `$XDG_DATA_HOME/QectorWorkbench` or `~/.local/share/QectorWorkbench`.

12. Troubleshooting

The app does not appear after launch.	A splash screen appears within about a second and stays up until the main window is ready. The first launch is the slow one: it extracts the bundled decoder wheel, then loads a compiled extension. If nothing appears at all, read <code>logs/boot.log</code> in the per-user data directory; it records every provisioning step and the exact import error.
A decoder is refused for a qLDPC code.	Choose <code>bp_osd</code> , <code>blossom</code> , <code>sparse_blossom</code> , <code>hybrid</code> , <code>predecoded</code> , <code>auto</code> or <code>auto_router</code> , or let the resilient path pick a working decoder.
CUDA reports unavailable.	Expected without a supported NVIDIA GPU and a healthy driver. The CPU backend decodes fully.
OpenCL always reports unavailable, even with a working OpenCL GPU.	Expected. The published <code>qector-decoder-v3</code> wheel is built without its OpenCL feature, so the backend does not exist in that build; this is not a driver or configuration fault and no environment variable enables it. The Hardware tab and 'qector hardware' show how many OpenCL devices the host itself exposes, so you can tell the two situations apart. Use CUDA or CPU.
Benchmarks differ between machines.	Throughput, latency and logical error rate are simulation outputs; they depend on hardware, seed and workload. Regenerate them for your setting.

13. Licensing and support

The software is source-available. Academic, personal and non-commercial research use is free. Commercial use requires a paid licence, and a 60-day commercial evaluation is available and creditable against a licence. The qector-decoder-v3 backend is licensed separately. Consult the EULA shipped with the application for full terms.

Buy a licence: <https://qector.store/pricing>. The application also has a Buy Licence button in the Documentation tab, under Developer and Licensing, alongside the maintainer and contact details.

Sales and licensing enquiries: contact@qector.store.

Project home and support: <https://www.qector.store>. Attribution: Guillaume Lessard / iD01t Productions (ORCID 0009-0000-3465-3753).

14. Glossary

Stabiliser code	A quantum error-correcting code defined by a set of commuting parity checks (stabilisers) whose measurement reveals errors without disturbing the encoded information.
Syndrome	The pattern of parity-check outcomes. A decoder maps a syndrome to a correction that explains it.
Decoder	An algorithm that turns a syndrome into a correction. QECTOR ships ten.
Code distance	The size of the smallest error that can cause an undetected logical failure. Larger distance means stronger protection.
Graphlike code	A code whose errors can be represented as edges in a graph, so matching decoders (union-find, blossom) apply directly.
qLDPC code	A quantum low-density parity-check code (here bicycle and bivariate bicycle). Not graphlike; decode with BP-OSD or the routed decoders.
Logical error rate	The simulated probability that a decode leaves an undetected logical error. A simulation output, not a fixed benchmark.
MCP	Model Context Protocol. A JSON-RPC interface that lets automation and language-model agents drive the workbench headlessly.

Appendix A. Full MCP tool list

All 56 tools registered by the server, alphabetical.

ambiguity_cluster_decode	Decode using AmbiguityClusterDecoder for high noise or non-graphlike codes
analyze_code_family	Analyze a code family with an example code instance
batch_decode	Batch-decode sampled syndromes on cpu/cuda/opencl via backend.run_batch_decode
batch_decode_gpu	Batch-decode on an explicit compute backend (cpu/cuda/opencl) with honest availability reporting , unavailable GPU backends return status='unavailable' with a reason, never fake results
belief_match_decode	Convenience seeded decode pinned to the belief_matching kind (BP posteriors + exact MWPM with faithfulness fallback)
benchmark_decoder	Benchmark a decoder on a code family via backend.run_benchmark (latency percentiles, throughput, logical error rate)
check_updates	Report whether the installed decoder backend matches the bundled release baseline (offline , no update service)
clear_decoder_cache	Clear the backend's native decoder cache
clear_results	Clear all stored benchmark results
colour_code_decode	Decode color code using BP-OSD over undecomposed detector error model
compare_benchmarks	Compare stored benchmark results side by side (throughput, p99 latency, logical error rate)
compat_report	Report ecosystem-integration availability (stim/sinter/pymatching/qiskit/ldpc) and research components
compatible_decoders	Live probe: which decoder kinds construct and produce a syndrome-verified correction on this code
decode_single	Run one seeded decode and report correction weight, syndrome validity and logical failure
decode_syndrome	Decode an explicit 0/1 syndrome (length n_checks) with a chosen decoder; syndrome_valid is the GF(2) re-check, logical_failure is null (no reference error exists, so it is unknowable)
decode_with_options	Seeded decode with validated per-decoder construction options (bp_osd bp_method/osd_order, hybrid_cascade escalation, GNN architecture); reports options_applied honestly
delete_resource	Delete a resource by ID
diagnostic_decode	Rich single decode via the backend's native decode_with_diagnostics (matched weight, backend used, internal fallback, timing, logicals)
doctor_diagnostics	Run system health and environment diagnostic checks via qd.doctor
export_benchmark	Export a stored benchmark result (by result_id) to the export directory
flush_usage	Flush usage metrics to Stripe metered billing API
generate_documentation	Generate code documentation files
get_code_properties	Get properties of a code family
get_config	Get current server configuration
get_decoder_info	Get information about a decoder

get_hardware_info	Get hardware/backend availability
get_resource	Get a specific resource by ID
get_resources	List all resources
get_results	Get stored benchmark results (most recent first-in order)
get_statistics	Get server statistics
get_system_info	Get system information
gnn_belief_match_decode	Convenience seeded decode pinned to the gnn_belief_matching kind with optional GNN architecture overrides (gnn_hidden_size, gnn_n_layers)
hybrid_cascade_stats	Batch-decode through the hybrid_cascade decoder and expose its live cascade statistics (prefilter_hits, escalations, hit rate, throughput, syndrome-match rate, logical error rate)
list_clients	List registered clients
list_code_families	List available code families
list_codes	List workbench code families plus the backend's native code catalogue
list_decoders	List available decoders
list_tools	List all available MCP tools
mcp_status	Get MCP server status
native_recommend	Backend-native decoder recommendation (qector_decoder_v3.recommend) with the mapped workbench decoder_kind
native_streaming	Native hardware-accelerated sliding-window streaming decode (qector_decoder_v3.sliding_window_decode) with per-round validity + telemetry
neural_predecoder_train	Train the NeuralPredecoder research/lab MLP on seeded (syndrome, error) pairs and evaluate on a disjoint held-out stream (exact-match, bit accuracy, syndrome validity, LER)
probe_decoders	Probe which decoders produce a valid (syndrome-verified) correction for a code , a self-test across every wired decoder
recommend_decoder	Recommend a decoder for a code/priority using detected hardware
register_client	Register a client
reset_config	Reset configuration to defaults
resilient_decode	Single decode with automatic multi-decoder fallback and a full attempt trace (autodebug.resilient_single_decode)
run_benchmark	Run a benchmark and store the result under a generated result_id
self_diagnostics	Run a full environment/decoder/hardware self-diagnostics report (autodebug.run_self_diagnostics)
set_config	Merge key/value pairs into the server configuration
set_license_key_file	Set license key file path for offline verification
sparse_blossom_radix_neighbors	Discover k-nearest candidate edges via SparseBlossom RadixHeap
stream_decode	Run a sliding-window streaming decode session via backend.run_streaming_session
two_stage_decode	Decode using TwoStageDecoder (decoupled X/Z sector decoders)
verify_license_token	Verify an Ed25519 signed license token string

version_info	App + decoder-backend version report (workbench baseline + installed backend, resolved locally , no network)